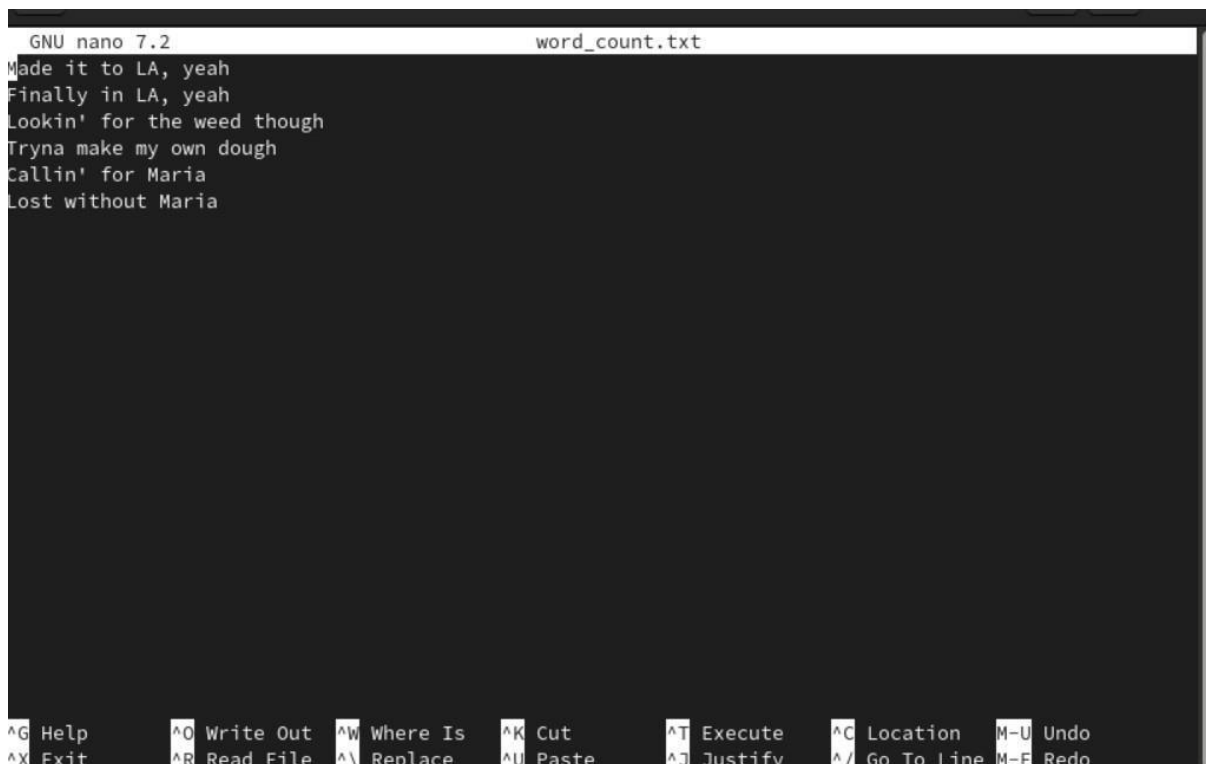


**Exp No: 2****Run a basic Word Count Map Reduce program to understand Map Reduce Paradigm****Aim:**

To Run a basic Word Count MapReduce program to understand Map Reduce Paradigm.

**Procedure:****Step 1: Create Data File:**

Create a file named "word\_count\_data.txt" and populate it with text data that you wish to analyze. Login with your Hadoop user.



```
GNU nano 7.2 word_count.txt
Made it to LA, yeah
Finally in LA, yeah
Lookin' for the weed though
Tryna make my own dough
Callin' for Maria
Lost without Maria

^G Help      ^O Write Out  ^W Where Is   ^K Cut        ^T Execute    ^G Location   M-U Undo
^X Exit      ^R Read File  ^\ Replace    ^U Paste      ^J Justify    ^_ Go To Line  M-E Redo
```

**Step 2: Mapper Logic - mapper.py:**

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

```
nano mapper.py
```

```
# Copy and paste the mapper.py code
```

```
#!/usr/bin/env python3
```

```
# import sys because we need to read and write data to STDIN and STDOUT
```

```
#!/usr/bin/python3
```

```
import sys for line
```

```
in sys.stdin:
```

```
    line = line.strip()
```

```
    # remove leading and trailing whitespace words
```

```
    = line.split()
```

```
# split the line into words for word in words: nano
word_count.txt print( '%s\t%s' % (word, 1)) Step
```

### 3: Reducer Logic - reducer.py:

Create a file named "reducer.py" to implement the logic for the reducer. The reducer will aggregate the occurrences of each word and generate the final output.

```
nano reducer.py
```

```
# Copy and paste the reducer.py code reducer.py
```

```
#!/usr/bin/python3      from
operator import itemgetter
import sys current_word =
None current_count = 0 word =
None for line in sys.stdin: line
= line.strip()
    word, count = line.split('\t', 1) try:
        count = int(count)
    except ValueError: continue
    if current_word == word:
        current_count += count else:
        if current_word: print( '%s\t%s' % (current_word,
            current_count))
        current_count = count current_word
        = word
    if current_word == word:
        print( '%s\t%s' % (current_word, current_count))
```

### Step 4: Prepare Hadoop Environment:

Start the Hadoop daemons and create a directory in HDFS to store your data. start-all.sh

```
hdfsdfs -mkdir /word_count_in_python
```

```
hdfsdfs -copyFromLocal /path/to/word_count.txt/word_count_in_python
```

### Step 5: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.

```
chmod 777 mapper.py reducer.py
```

### Step 6: Run Word Count using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access.

Then run the Word Count program using Hadoop Streaming.

```
hadoop jar /path/to/hadoop-streaming-3.3.6.jar \
```

```
-input /word_count_in_python/word_count_data.txt \
```

```
-output /word_count_in_python/new_output \
```

-mapper /path/to/mapper.py \

-reducer /path/to/reducer.py

```
C:\hadoop\sbin>hadoop jar C:\hadoop\share\hadoop\tools\lib\hadoop-streaming-3.3.6.jar ^
More? -input /user/hadoop/input/data.txt ^
More? -output /user/output ^
More? -mapper "python C:\Users\nithu\OneDrive\Documents\wordcount\mapper.py" ^
More? -reducer "python C:\Users\nithu\OneDrive\Documents\wordcount\reducer.py"
packageJobJar: [/C:/Users/nithu/AppData/Local/Temp/hadoop-unjar4804848770360266759/] [] C:\Users\nithu\AppData\Local\Temp\streamjob1651486068095611045.jar tmpDir=null
2024-09-14 21:53:11,332 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2024-09-14 21:53:11,629 INFO client.DefaultNoHARMAFailoverProxyProvider: Connecting to ResourceManager at /0.0.0.0:8032
2024-09-14 21:53:17,672 INFO mapreduce.JobResourceUploader: Disabling Erasure Coding for path: /tmp/hadoop-yarn/staging/nithu/.staging/job_1726328178557_0001
2024-09-14 21:53:18,139 INFO mapred.FileInputFormat: Total input files to process : 1
2024-09-14 21:53:18,244 INFO mapreduce.JobSubmitter: number of splits:2
2024-09-14 21:53:18,477 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1726328178557_0001
2024-09-14 21:53:18,477 INFO mapreduce.JobSubmitter: Executing with tokens: []
2024-09-14 21:53:18,675 INFO conf.Configuration: resource-types.xml not found
2024-09-14 21:53:18,676 INFO resource.ResourceUtils: Unable to find 'resource-types.xml'.
2024-09-14 21:53:19,371 INFO impl.YarnClientImpl: Submitted application application_1726328178557_0001
2024-09-14 21:53:19,428 INFO mapreduce.Job: The url to track the job: http://Nithisha:8088/proxy/application_1726328178557_0001/
2024-09-14 21:53:19,430 INFO mapreduce.Job: Running job: job_1726328178557_0001
2024-09-14 21:53:40,775 INFO mapreduce.Job: Job job_1726328178557_0001 running in uber mode : false
2024-09-14 21:53:40,781 INFO mapreduce.Job: map 0% reduce 0%
2024-09-14 21:53:45,938 INFO mapreduce.Job: map 50% reduce 0%
2024-09-14 21:53:45,951 INFO mapreduce.Job: map 100% reduce 0%
```

## Step 8: Check Output:

Check the output of the Word Count program in the specified HDFS output directory.

hdfs dfs -cat /word\_count\_in\_python/new\_output/part-00000

```
osboxes@fedora:~$ hdfs dfs -cat /output/part-r-00000
all 1
and 3
at 2
daylight 1
drank 1
from 2
hate 2
hiding 1
i 5
it 4
love 2
of 1
oh 2
our 1
poison 1
same 3
sins 1
the 5
time 2
wine 1
you 1
osboxes@fedora:~$
```

**Result:**

Thus, the program for basic Word Count Map Reduce has been executed successfully.