

description		dataset.txt																
File	Edit	View																
23907	20150101	2.423	-98.08	30.62	2.2	-0.6	0.8	0.9	7.0	1.47	C	3.7	1.1	2.5	99.9	85.4		
97.2	0.369	0.308	-99.000	-99.000	-99.000	7.0	8.1	-9999.0	-9999.0	-9999.0								
23907	20150102	2.423	-98.08	30.62	3.5	1.3	2.4	2.2	10.2	1.43	C	4.9	2.3	3.1	100.0	98.8		
99.8	0.391	0.327	-99.000	-99.000	-99.000	7.1	7.9	-9999.0	-9999.0	-9999.0								
23907	20150103	2.423	-98.08	30.62	15.9	2.3	9.1	7.5	3.1	11.00	C	16.4	2.9	7.3	100.0	34.8		
73.7	0.450	0.397	-99.000	-99.000	-99.000	7.6	7.9	-9999.0	-9999.0	-9999.0								
23907	20150104	2.423	-98.08	30.62	9.2	-1.3	3.9	4.2	0.0	13.24	C	12.4	-0.5	4.9	82.0	40.6		
61.7	0.413	0.352	-99.000	-99.000	-99.000	7.3	7.9	-9999.0	-9999.0	-9999.0								
23907	20150105	2.423	-98.08	30.62	10.9	-3.7	3.6	2.6	0.0	13.37	C	14.7	-3.0	3.8	77.9	33.3		
57.4	0.399	0.340	-99.000	-99.000	-99.000	6.3	7.0	-9999.0	-9999.0	-9999.0								
23907	20150106	2.423	-98.08	30.62	20.2	2.9	11.6	10.9	0.0	12.90	C	22.0	1.6	9.9	67.7	30.2		
49.3	0.395	0.335	-99.000	-99.000	-99.000	8.0	8.0	-9999.0	-9999.0	-9999.0								
23907	20150107	2.423	-98.08	30.62	10.9	-3.4	3.8	4.5	0.0	12.68	C	12.4	-2.1	5.5	82.7	36.5		
55.7	0.387	0.328	-99.000	-99.000	-99.000	7.6	8.3	-9999.0	-9999.0	-9999.0								
23907	20150108	2.423	-98.08	30.62	0.6	-7.9	-3.6	-3.3	0.0	4.98	C	3.9	-4.8	-0.5	57.7	37.6		
48.1	0.372	0.316	-99.000	-99.000	-99.000	4.7	6.1	-9999.0	-9999.0	-9999.0								
23907	20150109	2.423	-98.08	30.62	2.0	0.1	1.0	0.8	0.0	2.52	C	4.1	1.2	2.5	87.8	48.9		
64.4	0.368	0.312	-99.000	-99.000	-99.000	5.4	6.2	-9999.0	-9999.0	-9999.0								
23907	20150110	2.423	-98.08	30.62	0.5	-2.0	-0.8	-0.6	3.9	2.11	C	2.5	-0.1	1.4	99.9	47.7		
85.8	0.373	0.314	-99.000	-99.000	-99.000	5.1	6.0	-9999.0	-9999.0	-9999.0								
23907	20150111	2.423	-98.08	30.62	10.9	0.0	5.4	4.4	2.6	6.38	C	12.7	1.3	5.8	100.0	77.8		
97.1	0.420	0.362	-99.000	-99.000	-99.000	6.5	6.7	-9999.0	-9999.0	-9999.0								
23907	20150112	2.423	-98.08	30.62	6.5	1.4	4.0	4.3	0.0	1.55	C	6.9	2.7	5.1	100.0	89.4		
97.8	0.412	0.350	-99.000	-99.000	-99.000	7.3	7.5	-9999.0	-9999.0	-9999.0								
23907	20150113	2.423	-98.08	30.62	3.0	-0.7	1.1	1.2	0.0	3.26	C	5.6	0.7	2.9	99.7	80.7		
90.7	0.401	0.337	-99.000	-99.000	-99.000	6.1	6.8	-9999.0	-9999.0	-9999.0								
23907	20150114	2.423	-98.08	30.62	2.9	0.9	1.9	1.8	0.7	1.88	C	4.7	2.0	3.1	99.6	90.8		
97.9	0.395	0.331	-99.000	-99.000	-99.000	6.1	6.7	-9999.0	-9999.0	-9999.0								
23907	20150115	2.423	-98.08	30.62	13.2	1.2	7.2	6.4	0.0	13.37	C	16.4	1.4	6.7	98.9	46.7		
73.4	0.395	0.333	-99.000	-99.000	-99.000	6.7	7.0	-9999.0	-9999.0	-9999.0								
23907	20150116	2.423	-98.08	30.62	16.7	3.5	10.1	9.9	0.0	13.68	C	19.2	1.3	8.7	80.2	38.1		
58.2	0.391	0.330	-99.000	-99.000	-99.000	7.3	7.4	-9999.0	-9999.0	-9999.0								
23907	20150117	2.423	-98.08	30.62	19.5	5.0	12.2	12.3	0.0	10.96	C	20.9	3.3	10.6	87.7	30.4		
55.7	0.388	0.327	-99.000	-99.000	-99.000	8.7	8.4	-9999.0	-9999.0	-9999.0								
23907	20150118	2.423	-98.08	30.62	20.9	7.6	14.3	13.7	0.0	15.03	C	23.4	3.5	11.9	45.9	14.6		
31.4	0.383	0.325	-99.000	-99.000	-99.000	9.5	9.2	-9999.0	-9999.0	-9999.0								

Exp No 3

Map Reduce program to process a weather dataset

Aim:

To implement MapReduce program to process a weather dataset.

Procedure:

Step 1: Create Data File:

Create a file named "word_count_data.txt" and populate it with text data that you wish to analyse.

Login with your hadoop user.

Step 2: Mapper Logic - mapper.py:

Create a file named "mapper.py" to implement the logic for the mapper. The mapper will read input data from STDIN, split lines into words, and output each word with its count.

```
nano mapper.py
```

```
# Copy and paste the mapper.py code
```

```
#!/usr/bin/env python
```

```

import sys

# input comes from STDIN (standard input)
# the mapper will get daily max temperature and group it by month. so output will be
(month,daily_max_temperature)

Download the dataset (weather data) for
line in sys.stdin:

# remove leading and trailing whitespace line
= line.strip()

# split the line into words words
= line.split()

#See the README hosted on the weather website which help us understand how each
position represents a column
month      =      line[10:12]
daily_max  =      line[38:45]
daily_max          =
daily_max.strip()

# increase counters for
word in words:

# write the results to STDOUT (standard output);
# what we output here will be go through the shuffle process and then
# be the input for the Reduce step, i.e. the input for reducer.py #
# tab-delimited; month and daily max temperature as output print
('%s\t%s' % (month ,daily_max))

.

```

Step 3: Reducer Logic - reducer.py:

Create a file named "reducer.py" to implement the logic for the reducer. The reducer will aggregate the occurrences of each word and generate the final output.

nano reducer.py

```

# Copy and paste the reducer.py code
reducer.py

```



```
#!/usr/bin/env python from
operator import itemgetter
import sys

#reducer will get the input from stdid which will be a collection of key, value(Key=month ,
value= daily max temperature)

#reducer logic: will get all the daily max temperature for a month and find max temperature
for the month

#shuffle will ensure that key are sorted(month)
current_month = None current_max = 0 month
= None

# input comes from STDIN for
line in sys.stdin:

# remove leading and trailing whitespace line
= line.strip()

# parse the input we got from mapper.py month,
daily_max = line.split('\t', 1)

# convert daily_max (currently a string) to float
try:
daily_max = float(daily_max) except
ValueError:
# daily_max was not a number, so silently
# ignore/discard this line
continue

# this IF-switch only works because Hadoop shuffle process sorts map output
# by key (here: month) before it is passed to the reducer
if current_month == month: if daily_max >
current_max: current_max = daily_max
```



```

else: if
current_month:
# write result to STDOUT
print ("%s\t%s" % (current_month, current_max))
current_max = daily_max current_month =
month # output of the last month if current_month
== month:
print ("%s\t%s" % (current_month, current_max)) Step

```

4: Prepare Hadoop Environment:

Start the Hadoop daemons and create a directory in HDFS to store your data. start-all.sh

Step 5: Make Python Files Executable:

Give executable permissions to your mapper.py and reducer.py files.
 chmod 777 mapper.py reducer.py

Step 6: Run the program using Hadoop Streaming:

Download the latest hadoop-streaming jar file and place it in a location you can easily access.

Then run the program using Hadoop Streaming.

```

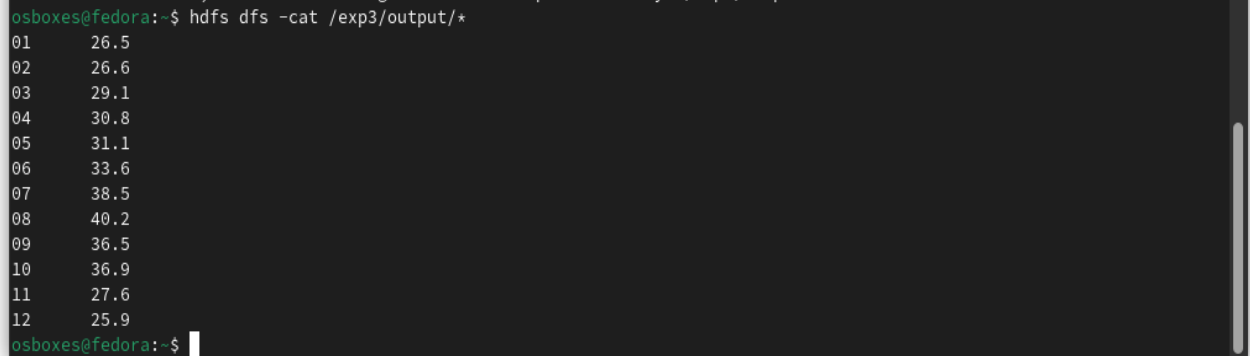
hadoop fs -mkdir -p /weatherdata
hadoop fs -copyFromLocal /home/sx/Downloads/dataset.txt /weatherdata hdfs
dfs -ls /weatherdata
hadoop jar /home/sx/hadoop-3.2.3/share/hadoop/tools/lib/hadoop-streaming-3.2.3.jar \
-input /weatherdata/dataset.txt \
-output /weatherdata/output \
-file "/home/sx/Downloads/mapper.py" \
-mapper "python3 mapper.py" \
-file "/home/sx/Downloads/reducer.py" \
-reducer "python3 reducer.py" hdfs dfs -text /weatherdata/output/* >
/home/sx/Downloads/outputfile.txt

```


Step 7: Check Output:

Check the output of the program in the specified HDFS output directory.

```
hdfs dfs -text /weatherdata/output/* > /home/sx/Downloads/output/  
/part-00000
```

A terminal window with a dark background. The prompt is 'osboxes@fedora:~\$'. The command 'hdfs dfs -cat /exp3/output/*' has been executed. The output is a list of 12 lines, each with an index and a temperature value. The prompt is now 'osboxes@fedora:~\$' with a cursor.

```
osboxes@fedora:~$ hdfs dfs -cat /exp3/output/*  
01      26.5  
02      26.6  
03      29.1  
04      30.8  
05      31.1  
06      33.6  
07      38.5  
08      40.2  
09      36.5  
10      36.9  
11      27.6  
12      25.9  
osboxes@fedora:~$
```

After copy and paste the above output in your local file give the below command to

remove the directory from hdfs :

```
hadoop fs -rm -r /weatherdata/output
```

Result:

Thus, the program for weather dataset using Map Reduce has been executed successfully.